

Explanation of Modifications in system calls

This document will explain the changes done in the kernel source files along with their snippets.

For open() and close() system call:

The code snippet below shows the myopen() in open.c file in the linux-4.19.281/fs directory. . It also displays additional information using printk to the kernel logs when a file is opened.

```
long do_sys_myopen(int dfd, const char __user *filename, int flags, umode_t mode)
{
    struct open_flags op;
    int fd = build_open_flags(flags, mode, &op);
    struct filename *tmp;

    if (fd)
        return fd;

    tmp = getname(filename);
    if (IS_ERR(tmp))
        return PTR_ERR(tmp);

    fd = get_unused_fd_flags(flags);
    if (fd >= 0) {
        struct file *f = do_filp_open(dfd, tmp, &op);
        if (IS_ERR(f)) {
            put_unused_fd(fd);
            fd = PTR_ERR(f);
        } else {
            fsnotify_open(f);
            fd_install(fd, f);
        }
    }
    putname(tmp);
    return fd;
}

SYSCALL_DEFINE3(myopen, const char __user *, filename, int, flags, umode_t, mode)
{
    printk("This is modified open() named as myopen(). This is done bu Afza Fatima,Nida Fatima and Soha Junaid.");
    printk("Filename: %s\n", filename);
    if (force_o_largefile())
        flags |= O_LARGEFILE;

    return do_sys_myopen(AT_FDCWD, filename, flags, mode);
}
```

The code snippet below shows the myclose() in file.c file in the linux-4.19.281/fs directory.

```

*file.c
/kernels/linux-4.19.281/fs

out_unlock:
    spin_unlock(&files->file_lock);
    return -EBADF;
}
EXPORT_SYMBOL(__close_fd); /* for ksys_close() */

int __myclose_fd(struct files_struct *files, unsigned fd)
{
    struct file *file;
    struct fdtable *fdt;

    spin_lock(&files->file_lock);
    fdt = files_fdtable(files);
    if (fd >= fdt->max_fds)
        goto out_unlock;
    fd = array_index_nospec(fd, fdt->max_fds);
    file = fdt->fd[fd];
    if (!file)
        goto out_unlock;
    rcu_assign_pointer(fdt->fd[fd], NULL);
    __put_unused_fd(files, fd);
    spin_unlock(&files->file_lock);
    return filp_close(file, files);

out_unlock:
    spin_unlock(&files->file_lock);
    return -EBADF;
}
EXPORT_SYMBOL(__myclose_fd); /* for ksys_close() */

void do_close_on_exec(struct files_struct *files)

```

Its system define can be found in open.c which prints additional information using printk in kernel logs.

```

SYSCALL_DEFINE1(myclose, unsigned int, fd)
{
    int retval = __myclose_fd(current->files, fd);
    printk("This is modified close() named as myclose() as part of OS PROJECT");
    /* can't restart close syscall because file table entry was cleared */
    if (unlikely(retval == -ERESTARTSYS ||
                 retval == -ERESTARTNOINTR ||
                 retval == -ERESTARTNOHAND ||
                 retval == -ERESTART_RESTARTBLOCK))
        retval = -EINTR;

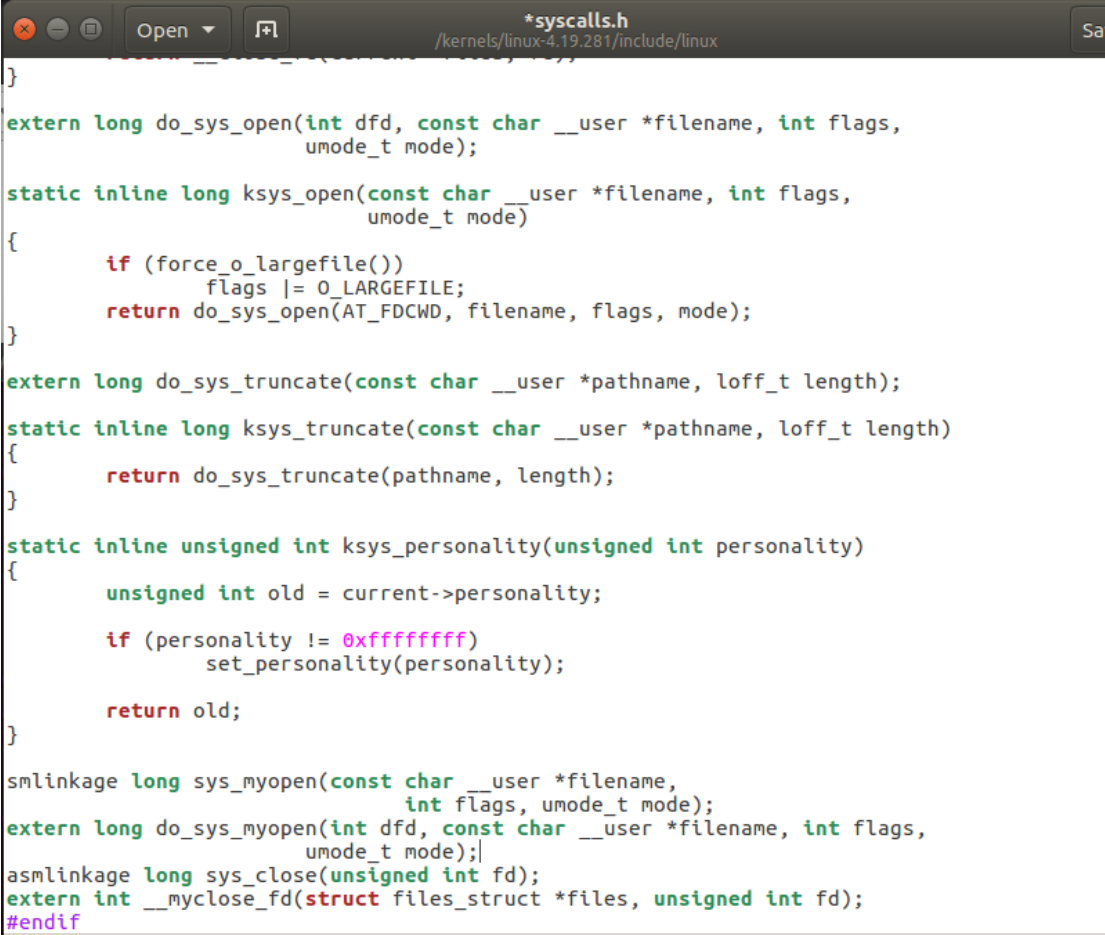
    return retval;
}

```

We add the modified calls in the syscall_64.tbl in the /arch/x86/entry/syscalls directory so that we can use the 335 and 336 reference number to make a call to this modified open and close call when we need to open or close the file.

324	common	membarrier	__x64_sys_membarrier
325	common	mlock2	__x64_sys_mlock2
326	common	copy_file_range	__x64_sys_copy_file_range
327	64	preadv2	__x64_sys_preadv2
328	64	pwritev2	__x64_sys_pwritev2
329	common	pkey_mprotect	__x64_sys_pkey_mprotect
330	common	pkey_alloc	__x64_sys_pkey_alloc
331	common	pkey_free	__x64_sys_pkey_free
332	common	statx	__x64_sys_statx
333	common	io_pgetevents	__x64_sys_io_pgetevents
334	common	rseq	__x64_sys_rseq
335	common	myopen	__x64_sys_myopen
336	common	myclose	__x64_sys_myclose

Then we write call declarations in the syscall.h header file in the directory /include/linux so that the kernel registers them while building



```

}

extern long do_sys_open(int dfd, const char __user *filename, int flags,
                       umode_t mode);

static inline long ksys_open(const char __user *filename, int flags,
                             umode_t mode)
{
    if (force_o_largefile())
        flags |= O_LARGEFILE;
    return do_sys_open(AT_FDCWD, filename, flags, mode);
}

extern long do_sys_truncate(const char __user *pathname, loff_t length);
static inline long ksys_truncate(const char __user *pathname, loff_t length)
{
    return do_sys_truncate(pathname, length);
}

static inline unsigned int ksys_personality(unsigned int personality)
{
    unsigned int old = current->personality;

    if (personality != 0xffffffff)
        set_personality(personality);

    return old;
}

asmlinkage long sys_myopen(const char __user *filename,
                           int flags, umode_t mode);
extern long do_sys_myopen(int dfd, const char __user *filename, int flags,
                          umode_t mode);
asmlinkage long sys_close(unsigned int fd);
extern int __myclose_fd(struct files_struct *files, unsigned int fd);
#endif

```

then we reboot the kernel after running make <target> and installing modules. That's how a system call is modified.